

lecture14a-neural-networks

March 11, 2022

1 Lecture 13: Neural Networks

1.0.1 Applied Machine Learning

Volodymyr Kuleshov, Jin Sun Cornell Tech

2 Part 1: An Artificial Neuron

In this lecture, we will learn about a new class of machine learning algorithms inspired by the brain. We will start by defining a few building blocks for these algorithms, and draw connections to neuroscience.

3 Review: Components of A Supervised Machine Learning Problem

At a high level, a supervised machine learning problem has the following structure:

$$\underbrace{\text{Training Dataset}}_{\text{Attributes + Features}} + \underbrace{\text{Learning Algorithm}}_{\text{Model Class + Objective + Optimizer}} \rightarrow \text{Predictive Model}$$

Where does the dataset come from?

4 Review: Binary Classification

In supervised learning, we fit a model of the form

$$f : \mathcal{X} \rightarrow \mathcal{Y}$$

that maps inputs $x \in \mathcal{X}$ to targets $y \in \mathcal{Y}$.

In classification, the space of targets \mathcal{Y} is *discrete*. Classification is binary if $\mathcal{Y} = \{0, 1\}$

Each value of y value is a *class* and we are interested in finding a hyperplane that separates the different classes.

5 Review: Logistic Regression

Logistic regression fits a model of the form

$$f(x) = \sigma(\theta^\top x) = \frac{1}{1 + \exp(-\theta^\top x)},$$

where

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$

is known as the *sigmoid* or *logistic* function.

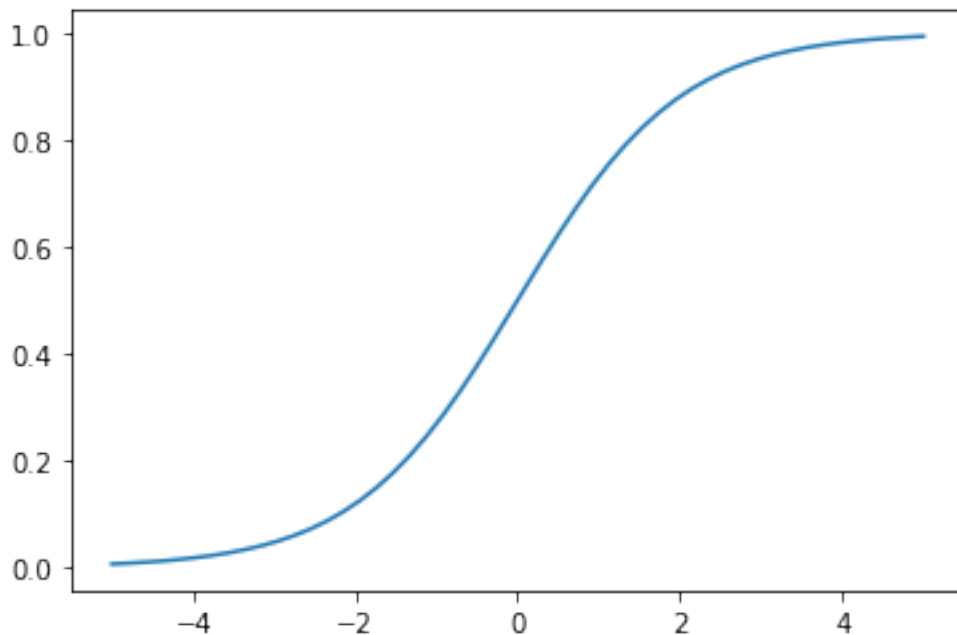
Here is how the logistic function looks like.

```
[1]: import numpy as np
      from matplotlib import pyplot as plt

      z = np.linspace(-5, 5)
      sigma = 1/(1+np.exp(-z))

      plt.plot(z, sigma)
```

```
[1]: [<matplotlib.lines.Line2D at 0x1172c9160>]
```



6 A Biological Neuron

In order to define an artificial neuron, let's look first a biological one.

TODO: PUT NEURON IMAGE FROM HERE: <https://cs231n.github.io/neural-networks-1/>

- Each neuron receives input signals from its dendrites
- It produces output signals along its axon, which connects to the dendrites of other neurons.

7 An Artificial Neuron: Example

We can imitate this machinery using an idealized artificial neuron. * The neuron receives signals x_j at dendrites, which are modulated multiplicatively: $w_j \cdot x_j$. * The body of the neuron sums the modulated inputs: $\sum_{j=1}^d w_j \cdot x_j$. * These go into the activation function that produces an output.

TODO: PUT ARTIFICIAL NEURON IMAGE FROM HERE: <https://cs231n.github.io/neural-networks-1/>

8 An Artificial Neuron: Notation

More formally, we say that a neuron is a model $f : \mathbb{R}^d \rightarrow [0, 1]$, with the following components: * Inputs x_1, x_2, \dots, x_d , denoted by a vector x . * Weight vector $w \in \mathbb{R}^d$ that modulates input x as $w^\top x$. * An activation function $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ that computes the output $\sigma(w^\top x)$ of the neuron based on the sum of modulated features $w^\top x$.

9 Logistic Regression as an Artificial Neuron

Logistic regression is a model of the form

$$f(x) = \sigma(\theta^\top x) = \frac{1}{1 + \exp(-\theta^\top x)},$$

that can be interpreted as a neuron that uses the *sigmoid* as the activation function.

10 Perceptron

Another model of a neuron.

11 Example

Need to implement a small example. Can probably copy-paste implementation of LR from the LR slides.

12 Activation Functions

Let's list a few.

Part 2: Artificial Neural Networks

Let's now see how we can connect neurons into networks that form complex models that further mimic the brain.

13 Review: Artificial Neuron

We say that a neuron is a model $f : \mathbb{R}^d \rightarrow [0, 1]$, with the following components:

- * Inputs x_1, x_2, \dots, x_d , denoted by a vector x .
- * Weight vector $w \in \mathbb{R}^d$ that modulates input x as $w^\top x$.
- * An activation function $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ that computes the output $\sigma(w^\top x)$ of the neuron based on the sum of modulated features $w^\top x$.

14 Review: Logistic Regression as Neuron

Logistic regression is a model of the form

$$f(x) = \sigma(\theta^\top x) = \frac{1}{1 + \exp(-\theta^\top x)},$$

that can be interpreted as a neuron that uses the *sigmoid* as the activation function.

15 Neural Networks: Intuition

A neural network is a directed graph in which a node is a neuron that takes as input the outputs of the neurons that are connected to it.

TODO: Add an image here. Maybe layer image from here: <https://cs231n.github.io/neural-networks-1/> (It probably needs some annotations)

Networks are typically organized in layers.

16 Neural Networks: Layers

A neural network layer is a model $f : \mathbb{R}^d \rightarrow \mathbb{R}^p$ that applies p neurons in parallel to an input x .

$$f(x) = \begin{bmatrix} \sigma(w_1^\top x) \\ \sigma(w_2^\top x) \\ \vdots \\ \sigma(w_p^\top x) \end{bmatrix}.$$

where each w_k is the vector of weights for the k -th neuron. We refer to p as the *size* of the layer.

By combining the w_k into one matrix W , we can write in a more succinct vectorized form:

$$f(x) = \sigma(W \cdot x) = \begin{bmatrix} \sigma(w_1^\top x) \\ \sigma(w_2^\top x) \\ \vdots \\ \sigma(w_p^\top x) \end{bmatrix},$$

where $\sigma(W \cdot x)_k = \sigma(w_k^\top x)$ and $W_{kj} = (w_k)_j$.

17 Neural Networks: Notation

A neural network is a model $f : \mathbb{R} \rightarrow \mathbb{R}$ that consists of a composition of L neural network layers:

$$f(x) = f_L \circ f_{L-1} \circ \dots \circ f_1(x).$$

The final layer f_L has size one (assuming the neural net has one output); intermediary layers f_l can have any number of neurons.

The notation $f \circ g(x)$ denotes the composition $f(g(x))$ of functions

18 Example of a Neural Network

- Let's implement a small neural net in the same that we implemented logistic regression
- Then we just run it

19 Types of Neural Network Layers

There are many types of neural network layers that can exist. Here are a few: * Output layer: normally has one neuron and special activation function that depends on the problem * Input layer: normally, this is just the input vector x . * Hidden layer: Any layer between input and output. * Dense layer: A layer in which every input is connected to every neuron. * Convolutional layer: A layer in which the operation $w^\top x$ implements a mathematical [convolution](#). * Anything else?

20 Neuroscience Angle

Anything we should say here?

Part 3: Backpropagation

We have defined what is an artificial neural network.

Let's not see how we can train it.

21 Review: Neural Network Layers

A neural network layer is a model $f : \mathbb{R}^d \rightarrow \mathbb{R}^p$ that applies p neurons in parallel to an input x .

$$f(x) = \sigma(W \cdot x) = \begin{bmatrix} \sigma(w_1^\top x) \\ \sigma(w_2^\top x) \\ \vdots \\ \sigma(w_p^\top x) \end{bmatrix},$$

where each w_k is the vector of weights for the k -th neuron and $W_{kj} = (w_k)_j$. We refer to p as the *size* of the layer.

22 Review: Neural Networks

A neural network is a model $f : \mathbb{R} \rightarrow \mathbb{R}$ that consists of a composition of L neural network layers:

$$f(x) = f_L \circ f_{L-1} \circ \dots \circ f_1(x).$$

The final layer f_L has size one (assuming the neural net has one output); intermediary layers f_l can have any number of neurons.

The notation $f \circ g(x)$ denotes the composition $f(g(x))$ of functions

TODO: Add some kind of image from the previous part of the lecture

23 Review: The Gradient

The gradient $\nabla_{\theta} f$ further extends the derivative to multivariate functions $f : \mathbb{R}^d \rightarrow \mathbb{R}$, and is defined at a point θ_0 as

$$\nabla_{\theta} f(\theta_0) = \begin{bmatrix} \frac{\partial f(\theta_0)}{\partial \theta_1} \\ \frac{\partial f(\theta_0)}{\partial \theta_2} \\ \vdots \\ \frac{\partial f(\theta_0)}{\partial \theta_d} \end{bmatrix}.$$

In other words, the j -th entry of the vector $\nabla_{\theta} f(\theta_0)$ is the partial derivative $\frac{\partial f(\theta_0)}{\partial \theta_j}$ of f with respect to the j -th component of θ .

24 Review: Gradient Descent

If we want to optimize an objective $J(\theta)$, we start with an initial guess θ_0 for the parameters and repeat the following update until the function is no longer decreasing:

$$\theta_i := \theta_{i-1} - \alpha \cdot \nabla_{\theta} J(\theta_{i-1}).$$

As code, this method may look as follows:

```
theta, theta_prev = random_initialization()
while abs(J(theta) - J(theta_prev)) > conv_threshold:
    theta_prev = theta
    theta = theta_prev - step_size * gradient(theta_prev)
```

25 Backpropagation

How do we apply gradient descent to a neural network?

Explain backpropagation

26 Review: Chain Rule of Calculus

Probably will need to review this at some point.

27 Example

Let's implement backprop with the simple NN model we had earlier.

Part 4: Stochastic Gradient Descent

In practice, neural networks are often trained on very large datasets.

This requires a modification to the gradient descent algorithm that we have seen earlier.

Volodymyr will create this section

[]: